

Aura Velocity

The Temporal Intelligence Terminal

A Personal Intelligence System for Time, Energy, and Life Capital Allocation

Document Version v1.0	Product Stage Live — Growth / Experimental	Author End-to-End Product Strategy & AI-Native Development	Last Updated May 2026	Live URL existence-stream.web.app
--------------------------	---	---	--------------------------	--

Live Link

01 PURPOSE & SCOPE OF THIS DOCUMENT

This PRD documents the product thinking, feature decisions, success criteria, and strategic rationale behind Aura Velocity. It is written retrospectively — meaning the product has already been built and shipped — but reflects the actual reasoning that guided every major decision during development.

The document is structured for two audiences: product reviewers who want to evaluate PM craft and decision-making quality, and future collaborators or contributors who need context on the system architecture, behavioural model, and product decisions.

Scope of This PRD

This document covers the full v1.0 feature set of Aura Velocity: the Temporal Integrity Engine, Subtraction Engine (AI), Circadian Engine, Aura Score, Resource Flow, Personal CRM, Ghost Mode, and the Mechanical UI design system. It does not cover Fit 90 (a separate product in the same portfolio).

02 PROBLEM STATEMENT

The Core Problem: Entropy Without Accountability

High-performing individuals — founders, engineers, and ambitious builders — consistently report the same frustration: they end their days feeling busy but unable to explain where their time went. This is not a scheduling failure. It is a visibility and accountability failure.

The problem has four compounding dimensions, each of which existing productivity tools address in isolation — but none of which they solve together:

#	Problem Dimension	Root Cause	Gap in Existing Tools
1	The Silent Time Leak	No visibility into low-value, habitual activity ("Boondoggles")	Tools optimise for task completion, not task quality
2	Circadian Misalignment	Scheduling by clock time ignores biological energy peaks	No tool maps activity intensity to chronotype
3	Financial Fog	Static bank balance obscures true financial runway	Finance apps track transactions, not survival time

4	Relationship Decay	No system maintains social capital proactively	CRMs exist for business, not personal relationships
---	--------------------	--	---

Problem Validation

Initial validation came through personal usage and manual behavioural tracking experiments conducted during early exploration. The initial hypothesis — that tracking "Regret" rather than "Progress" would be more behaviourally effective — was tested personally over several weeks using a manual Excel sheet. The frustrations that emerged from that process (non-aesthetic UI, no intelligence layer, no biological alignment) directly shaped the first iteration of the product.

Initial Product Insight

"I looked at Google Calendar, Notion, Todoist, and various habit trackers. Every one of them had the same flaw: they told you what to add. None of them asked the more important question — what should you stop doing?"

03 GOALS & NON - GOALS

Product Goals

The product optimises primarily for behavioural awareness and decision quality rather than raw productivity volume. The following goals guided every feature decision. They are ordered by priority — if a decision conflicted with a higher-priority goal, the lower-priority goal was deprioritised.

Priority	Goal	How Success Is Measured
G1 — Primary	Reduce time wasted on low-value activities	Hours reclaimed per month via Subtraction Engine
G2 — Primary	Align work output with biological energy peaks	Aura Score trend (week-over-week improvement)
G3 — Secondary	Surface the gap between intention and actual behaviour	Ghost Mode Delta (shrinking over time)
G4 — Secondary	Provide a unified view of all life capital (time, money, relationships)	Daily active use across all four modules
G5 — Tertiary	Make time-logging feel consequential, not administrative	Mechanical Hits Per Day (log frequency)

Non-Goals (What This Product Deliberately Does Not Do)

Defining what a product will not do is as important as defining what it will. The following were explicitly out of scope for v1.0:

- Team or collaborative features — Aura Velocity v1.0 is a single-user behavioural OS. Team functionality is a v3 roadmap item.
- Passive tracking — the product requires deliberate manual logging. Automatic time tracking was considered and rejected because it removes the psychological cost of time awareness.
- Broad user accessibility — the design is intentionally high-density and high-cognitive-load. Accessibility optimisation for new or casual users is not a v1.0 goal.

- Social sharing or community features — the product is a decision environment, not a social platform.
- Replacing existing calendars — Aura Velocity augments decision-making. It is not designed to replace Google Calendar for meeting management.

04 USER PERSONAS & JOBS TO BE DONE

Primary Persona — "Vampire": The Late-Night Builder

Attribute	Detail
Who they are	Indie builder or senior engineer. Age 22–38. AI-native. Works non-linear hours.
Core motivation	Protect high-leverage focus hours. Treat time as the most finite, non-renewable resource.
Key frustration	Ends days feeling occupied but not productive. No visibility into where cognitive energy went.
Emotional driver	Fear of "drifting" — living a year on accidental habits rather than intentional decisions.
Relationship with tools	Has tried and abandoned most productivity apps. Finds them patronising, soft, or generic.

Primary JTBD

"When I feel my week slipping away into low-value meetings and distractions, I want to run a Subtraction Audit so that I can systematically reduce low-value activity and protect my high-leverage focus hours."

Secondary Persona — The Technical PM / "The Architect"

Attribute	Detail
Who they are	Product Manager, strategy consultant, or data-driven professional. Age 24–40.
Core motivation	Systematic, whole-life optimisation. Wants one interface for all resource allocation decisions.
Key frustration	App-hopping between calendar, finance, CRM, habit tracker. Fragmented visibility across life systems.
Emotional driver	Derives deep satisfaction from systems that work precisely and feedback that is honest.

Secondary JTBD

"When I feel overwhelmed by the complexity of managing time, money, and relationships across different apps, I want a single unified terminal so that I can see the full picture of my life capital at once and make better allocation decisions."

05 FEATURE REQUIREMENTS

Features are documented with their user problem, product requirement, acceptance criteria, and the key tradeoff that shaped the final decision. This structure reflects how each feature was actually reasoned through — not as a list of specifications, but as a series of deliberate product bets.

F1 — Subtraction Engine (AI Core)

Field	Detail
User Problem	Users accumulate low-value, high-regret activities without realising the cumulative cost.
Requirement	The system must analyse activity logs over a rolling 30-day period and generate specific behavioural proposals: Purge (eliminate entirely), Reduce (cut frequency or duration), or Merge (combine with a similar activity).
AI Model	Gemini 1.5 Flash. Receives a structured JSON payload of activity name, duration, regret rating, energy level, and chronotype alignment score. Returns a structured Subtraction Proposal. Proposal quality varies significantly based on logging consistency and the granularity of activity naming.
Acceptance Criteria	(1) Proposal generated after minimum 14 days of data. (2) Each proposal includes a specific activity name, recommended action, and rationale. (3) Local fallback displays last-generated proposal if API unavailable.
Key Tradeoff	Conversational AI (chatbot) vs. forensic auditor. Chose forensic auditor: higher output quality, predictable structure, no conversation maintenance overhead.

F2 — Temporal Integrity Engine

Field	Detail
User Problem	Overlapping or inconsistent time logs undermine the credibility of all downstream analytics.
Requirement	The engine must enforce temporal exclusivity at all times. No two activity blocks can occupy the same time slice. Conflicts must be resolved automatically using a greedy forward-push cascade — the later activity is pushed forward to the next available slot.
Custom Logic	temporalEngine.js — handles chronological sorting, exclusivity enforcement, and the Heal Engine (cross-device sync conflict resolution).
Acceptance Criteria	(1) Zero overlapping logs in any user session. (2) Heal Engine resolves conflicts silently without user intervention. (3) All edits preserve total logged hours within a given day.
Key Tradeoff	Manual conflict resolution (user decides) vs. automatic resolution. Chose automatic: cognitive load of conflict prompts defeats the product's speed goal.

F3 — Circadian Engine & Aura Score

Field	Detail
User Problem	Users schedule high-intensity tasks during low-energy windows, reducing both output quality and perceived productivity.
Requirement	The engine must map the user's self-selected chronotype (Lark, Owl, Balanced, Vampire) to a daily energy curve. Each logged activity's intensity level is compared against the energy curve at that time slot. The delta produces the Aura Score — an XP-based alignment metric.
Design Decision	Four chronotype categories intentionally simplify biological complexity. Usability and fast onboarding were prioritised over scientific precision.
Acceptance Criteria	(1) Aura Score updates in real time as activities are logged. (2) Score visibly differentiates between aligned and misaligned sessions. (3) Weekly trend view shows score improvement over time.
Key Tradeoff	Scientific precision (continuous energy curve) vs. usability (4 categories). Chose usability: over-engineering the biological model reduces adoption without improving outcomes.

F4 — Resource Flow (Finance Terminal)

Field	Detail
User Problem	A static bank balance creates Financial Fog — no visibility into how long current resources will last.
Requirement	The module must calculate and display: (1) Liquidity Pulse — total available capital with a real-time visual indicator. (2) Burn Rate — average daily expenditure over the current monthly cycle. (3) Runway — total days of financial freedom at current burn rate.
Product Framing	Deliberately avoids "savings" language. Reframes all finance as startup capital: spending is "Burn," balance is "Runway." Language is an intentional product decision.
Acceptance Criteria	(1) Runway updates with every new transaction. (2) Liquidity Pulse reflects real-time balance. (3) Burn Rate calculated on rolling 30-day average.
Key Tradeoff	Full budgeting system vs. survival metrics only. Chose survival metrics: the insight value of Runway far exceeds traditional budget categories for the target persona. Early usage suggested that financial tracking engagement varied significantly depending on whether users already tracked expenses elsewhere.

F5 — Personal CRM (Relationship Capital)

Field	Detail
User Problem	Relationships with mentors, collaborators, and close contacts decay silently because there is no proactive maintenance system.
Requirement	The module must: (1) Track last contact date per person. (2) Assign each contact a relationship tier (Close, Professional, Acquaintance). (3) Generate nudges when contact cadence falls below the tier threshold.
Acceptance Criteria	(1) Nudge triggered at the correct tier-specific interval. (2) Last contact date updates on manual log entry. (3) Tier reassignment updates nudge frequency immediately.
Key Tradeoff	Automated contact logging vs. manual entry. Chose manual: intentional contact reinforces the relationship, which is the product's point.

F6 — Ghost Mode

Field	Detail
User Problem	Users have no honest feedback loop between what they intend to do and what they actually do.
Requirement	Ghost Mode must allow users to pre-log a planned schedule ("Ghost") for a day or week. At end of day, actual logs are overlaid on the Ghost template. The Delta — the difference between planned and actual — is surfaced as the primary metric.
Design Philosophy	Designed for awareness, not punishment. The Ghost Mode delta is presented neutrally — as data, not as failure.
Acceptance Criteria	(1) Ghost templates creatable up to 7 days in advance. (2) Delta calculated and displayed per activity and per day. (3) Ghost templates reusable as recurring schedules.
Key Tradeoff	Real-time prompting (alert when deviating) vs. retrospective delta. Chose retrospective: real-time prompts create anxiety and interrupt focus.

F7 — Mechanical UI Design System

Field	Detail
User Problem	Digital time-logging feels weightless and therefore creates no psychological accountability.
Requirement	Every primary user action (log entry, activity completion, deletion) must trigger: (1) A distinct auditory "clunk" sound. (2) A micro-animation (screen pulse or shake). Living background animations must reflect the biological flow concept at all times.
Design Rationale	Deliberate friction is a feature, not a bug. Faster logging was explicitly deprioritised in favour of making each logging action feel consequential.
Acceptance Criteria	(1) Audio feedback fires within 50ms of action. (2) Animation completes within 300ms. (3) Mechanical feedback cannot be disabled (v1.0) — it is core to the product identity.
Key Tradeoff	Silent/minimal UI (standard SaaS) vs. Mechanical UI. Chose Mechanical UI: the tactile feedback loop is the primary retention mechanism for the target persona.

06 PRIORITISATION FRAMEWORK

Feature Priority Matrix

Features were prioritised using a simplified Impact vs. Effort framework, weighted against alignment with the North Star Metric (Hours Reclaimed Per Month). P0 = must ship for product to be viable. P1 = high value, ship in v1.0. P2 = valuable but deferrable. Features that generated reusable behavioural infrastructure were prioritised over isolated utility features.

Feature	Impact	Effort	Priority	Rationale
Temporal Integrity Engine	High	High	P0	Without temporal exclusivity, all downstream analytics are invalid. Product cannot exist without this.
Subtraction Engine (AI)	High	High	P0	Core product differentiator. Entire brand promise depends on this feature existing.
Circadian Engine + Aura Score	High	Medium	P0	Biological alignment is the second pillar of the product thesis. Required for v1.0.
Mechanical UI System	High	Medium	P0	Primary retention mechanism. Without tactile feedback, the product feels identical to any other logger.
Ghost Mode	High	Medium	P1	Highest-insight feature but requires sufficient log history. Ships in v1.0 as it deepens the engagement loop.
Resource Flow	Medium	Medium	P1	Adds Life Capital dimension. Not core to time tracking thesis but reinforces unified OS positioning.
Personal CRM	Medium	Low	P1	Low engineering lift, high philosophical value. Relationship capital completes the Life Capital model.
Strategy Forge / Blueprints	Medium	Medium	P2	Long-term planning layer. Valuable but not required to deliver the core value proposition.
Social Velocity (Team)	High	Very High	P2	V3 roadmap item. Significant scope expansion — deferred to avoid diluting v1.0 focus.

Prioritisation Rationale

The P0 cluster (Temporal Engine, Subtraction Engine, Circadian Engine, Mechanical UI) forms the minimum viable philosophy — the smallest set of features that makes Aura Velocity meaningfully different from any existing tool. Shipping without any one of these four would produce a product that fails to deliver on its core positioning. The sequencing also reflected solo-build execution constraints, requiring infrastructure-heavy systems to deliver multi-feature leverage.

The P1 cluster (Ghost Mode, Resource Flow, Personal CRM) expands the Life Capital model beyond time and energy. These features were included in v1.0 because they reinforce the unified OS narrative, but they could survive as v2 items without compromising core user value.

07 USER JOURNEY & KEY FLOWS

Core Engagement Loop

The product is designed around a daily engagement loop that compounds over time. Each day generates data. Data enables the weekly Aura Score review. After 30 days, data enables the Subtraction Audit — the product's highest-value moment.

Stage	Trigger	User Action	System Response	Value Delivered
Log	User completes an activity	Tap to log activity + duration + energy level	Temporal Engine validates, Aura Score updates, Mechanical clunk fires	Immediate feedback; time made visible
Rate	End of activity block	User assigns Regret Rating (1–10)	Rating stored; feeds Subtraction Engine dataset	Honest reflection on time quality
Review	End of day	User checks Ghost Mode delta	System shows planned vs. actual gap	Awareness of behavioural drift
Audit	Every 30 days	User triggers Subtraction Engine	Gemini analyses 30-day log; returns Purge/Reduce/Merge proposal	Hours reclaimed; core value delivered
Adjust	Post-audit	User acts on proposals	Activity library updated; future Aura Scores recalibrate	Compounding improvement over time

Onboarding Flow

Onboarding is designed as identity formation, not a feature walkthrough. The user is not asked to explore the app — they are asked to declare who they are.

1. Select Chronotype (Lark / Owl / Balanced / Vampire) — this is the first decision because it shapes the entire biological alignment layer.
2. Define primary activity categories and their intensity levels — the user builds their personal activity vocabulary.
3. Set up "Armor Equipping" tasks — the daily preparation rituals that bookend each session.
4. Log first activity — triggers the first Mechanical UI feedback and the first Aura Score data point.

Onboarding Philosophy

High friction onboarding is intentional. Every step requires a decision about identity and values. Users who complete onboarding have self-selected into the product's philosophy — they are significantly more likely to engage deeply and log consistently.

This approach trades top-of-funnel conversion for stronger downstream engagement quality.

08 TECHNICAL ARCHITECTURE — PRODUCT VIEW

This section describes the architecture not as an engineering specification, but as a series of product decisions implemented in code. Each architectural choice has a user-facing reason.

Architectural Decision	Technology	Product Reason	Trade-off Made
Dual-mirror state	LocalStorage + Firestore	Zero-latency interactions — the app feels native on the web	Sync complexity increased; the Heal Engine was required to manage conflicts
AI as batch auditor	Gemini 1.5 Flash (JSON prompts)	High-quality structured output; predictable and reliable proposal format	No real-time AI interaction; all analysis is asynchronous
Local AI fallbacks	Cached last-known proposals	Product functions fully without API access; AI is enrichment, not dependency	Proposals may be stale; risk is low given 30-day audit cadence
Monolithic context	AppContext.jsx (~2,500 lines)	Single context maximised build velocity during rapid AI-assisted iteration	Refactoring risk for team-based scaling; acceptable for v1.0
Custom Temporal Engine	temporalEngine.js	Domain logic too critical to trust to a generic library; exclusivity is the product's data integrity guarantee	Higher initial build time; offset by complete control over edge case handling

09 SUCCESS METRICS & MEASUREMENT PLAN

North Star Metric

Hours Reclaimed Per Month

The total time identified as low-value and removed from the user's schedule via the Subtraction Engine. This is the single number that defines whether the product is delivering on its core promise. Every other metric exists in service of this one.

The metric is directionally useful rather than scientifically precise, since reclaimed time depends partly on subjective user judgment.

Metric Framework

Metric	Category	Definition	Target Signal	Why It Matters
Hours Reclaimed / Month	North Star	Total hours pruned via Subtraction Engine proposals	Increasing month-over-month	Directly measures the product's core value delivery
Aura Score Trend	Engagement	Week-over-week change in circadian alignment score	Positive slope over 4 weeks	Measures whether the product is changing scheduling behaviour

Mechanical Hits / Day	Engagement	Number of log entries per day	Consistent 6–10 per day	Baseline signal for active use; too low = logging abandoned
Ghost Mode Delta	Behaviour Change	Gap between planned and actual schedule (hours)	Shrinking over 30 days	A shrinking delta proves the product is aligning intent with action
Subtraction Audit Rate	Retention	% of users who act on at least one Subtraction Proposal	Action taken on 1+ proposals per audit	Measures whether AI output translates to real behaviour change
Runway (Days)	Health	Days of financial freedom at current Burn Rate	Stable or increasing	Strategic life health metric; correlates with focus capacity

What Success Looks Like at 30 / 60 / 90 Days

Timeframe	Leading Indicator	Lagging Indicator
Day 30	User has completed their first Subtraction Audit	Aura Score shows positive trend; Ghost Mode delta < 2 hours/day
Day 60	User has acted on at least 2 Subtraction proposals	Measurable reduction in logged Boondoggle hours vs. Day 1
Day 90	User has reclaimed 20+ hours vs. their baseline	Runway (Resource Flow) stable or improved; CRM nudges being acted on

10 RISKS & MITIGATIONS

Risk	Category	Likelihood	Impact	Mitigation
High onboarding friction causes drop-off before activation	UX	High	High	Track completion rate per onboarding step; add optional "quick start" path that bypasses full setup in v1.1
AI Subtraction proposals feel generic or irrelevant	AI Quality	Medium	High	Improve JSON prompt engineering with more granular activity metadata; add user feedback mechanism on proposals (Accept / Reject / Modify)
Monolithic AppContext becomes unmaintainable	Technical	Medium	Medium	Acceptable for v1.0; refactor plan to modular contexts in v2.0 if team expands
Gemini API unavailability disrupts core feature	Infrastructure	Low	High	Local fallback already implemented — last-generated proposal surfaced; app fully functional without API
Users log inconsistently, making Subtraction Engine data insufficient	Behaviour	High	Medium	Mechanical UI + daily streak incentives designed to maintain log frequency; minimum 14-day data threshold before audit runs

Financial Fog module feels disconnected from time tracking	Product Coherence	Medium	Low	Roadmap item: connect Burn Rate to logged activity types (e.g. spending on Boondoggles reduces Aura Score)
--	-------------------	--------	-----	--

11 PRODUCT ROADMAP

The roadmap is sequenced around behavioural coherence — each phase deepens the Life Capital model rather than expanding feature surface area. New modules are added only when they reinforce the existing philosophical thesis.

Phase	Horizon	Focus	Key Features
v1.0 — Current	Live	Core Life Capital OS	Subtraction Engine, Temporal Engine, Circadian Engine, Aura Score, Resource Flow, Personal CRM, Ghost Mode, Mechanical UI
v1.1	Short-term (1–2 months)	Onboarding & AI quality	Quick-start onboarding path, Subtraction proposal feedback loop (Accept/Reject/Modify), Aura Score weekly digest notification
v2.0	Medium-term (3–6 months)	Expanded capital model	Learning & Growth module (books, courses, skill tracking), Fit 90 integration (unified biological performance view), Predictive Scheduling (AI suggests optimal time slots based on Aura history)
v2.5	Medium-term (6–9 months)	Intelligent forecasting	Predictive energy mapping from historical Aura Scores, Burn Rate projections, Relationship decay probability scoring
v3.0	Long-term (12+ months)	Social Velocity	Team Subtraction Audits for startups, Collective focus and burn rate dashboards, Organisational time leakage identification

12 OPEN QUESTIONS

The following questions remain open and should be addressed before v1.1 development begins. They are documented here because honest product thinking includes knowing what you don't yet know.

- Onboarding drop-off: At which step of the onboarding flow do users abandon most frequently? This data will determine whether the quick-start path is needed in v1.1 or can be deferred.
- Optimal audit cadence: Is 30 days the right interval for the Subtraction Engine? A shorter window (14 days) may produce more actionable proposals earlier in the user lifecycle.
- Regret Rating calibration: Users may rate activities inconsistently over time as their standards shift. Does the AI prompt need to anchor ratings to an explicit scale?
- Aura Score gamification depth: Is XP + alignment score sufficient motivation, or does the retention loop need a streak mechanic or social comparison element?

- Resource Flow adoption: Is the Finance Terminal being used by users who adopt the time-tracking modules? If not, it may be better surfaced as an opt-in module rather than a default tab.

Aura Velocity

Product Requirements Document · v1.0

existence-stream.web.app · React 19 + Firebase + Gemini AI